

# ML-RANDOMFORESTREGRESSION-CIHANERSOY

April 18, 2020

```
[1]: import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.metrics import mean_squared_error, accuracy_score
from sklearn.preprocessing import LabelEncoder
from scipy.stats import randint
```

Necessary libraries are downloaded. Pandas and numpy are standard data science libraries. Sklearn is for prediction and hyperparameter tuning. Scipy library is used to generate random numbers for RandomizedSearchCV.

```
[2]: df = pd.read_csv('DelayedFlights.csv')
```

DelayedFlights.csv dataset is brought into a pandas dataframe.

```
[3]: df.describe()
```

```
[3]:
```

	Unnamed: 0	Year	Month	DayofMonth	DayOfWeek	\
count	1.936758e+06	1936758.0	1.936758e+06	1.936758e+06	1.936758e+06	
mean	3.341651e+06	2008.0	6.111106e+00	1.575347e+01	3.984827e+00	
std	2.066065e+06	0.0	3.482546e+00	8.776272e+00	1.995966e+00	
min	0.000000e+00	2008.0	1.000000e+00	1.000000e+00	1.000000e+00	
25%	1.517452e+06	2008.0	3.000000e+00	8.000000e+00	2.000000e+00	
50%	3.242558e+06	2008.0	6.000000e+00	1.600000e+01	4.000000e+00	
75%	4.972467e+06	2008.0	9.000000e+00	2.300000e+01	6.000000e+00	
max	7.009727e+06	2008.0	1.200000e+01	3.100000e+01	7.000000e+00	

  

	DepTime	CRSDepTime	ArrTime	CRSArrTime	FlightNum	\
count	1.936758e+06	1.936758e+06	1.929648e+06	1.936758e+06	1.936758e+06	
mean	1.518534e+03	1.467473e+03	1.610141e+03	1.634225e+03	2.184263e+03	
std	4.504853e+02	4.247668e+02	5.481781e+02	4.646347e+02	1.944702e+03	
min	1.000000e+00	0.000000e+00	1.000000e+00	0.000000e+00	1.000000e+00	
25%	1.203000e+03	1.135000e+03	1.316000e+03	1.325000e+03	6.100000e+02	
50%	1.545000e+03	1.510000e+03	1.715000e+03	1.705000e+03	1.543000e+03	
75%	1.900000e+03	1.815000e+03	2.030000e+03	2.014000e+03	3.422000e+03	
max	2.400000e+03	2.359000e+03	2.400000e+03	2.400000e+03	9.742000e+03	

	...	Distance	TaxiIn	TaxiOut	Cancelled	\
count	...	1.936758e+06	1.929648e+06	1.936303e+06	1.936758e+06	
mean	...	7.656862e+02	6.812975e+00	1.823220e+01	3.268348e-04	
std	...	5.744797e+02	5.273595e+00	1.433853e+01	1.807562e-02	
min	...	1.100000e+01	0.000000e+00	0.000000e+00	0.000000e+00	
25%	...	3.380000e+02	4.000000e+00	1.000000e+01	0.000000e+00	
50%	...	6.060000e+02	6.000000e+00	1.400000e+01	0.000000e+00	
75%	...	9.980000e+02	8.000000e+00	2.100000e+01	0.000000e+00	
max	...	4.962000e+03	2.400000e+02	4.220000e+02	1.000000e+00	

  

	Diverted	CarrierDelay	WeatherDelay	NASDelay	SecurityDelay	\
count	1.936758e+06	1.247488e+06	1.247488e+06	1.247488e+06	1.247488e+06	
mean	4.003598e-03	1.917940e+01	3.703571e+00	1.502164e+01	9.013714e-02	
std	6.314722e-02	4.354621e+01	2.149290e+01	3.383305e+01	2.022714e+00	
min	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	
25%	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	
50%	0.000000e+00	2.000000e+00	0.000000e+00	2.000000e+00	0.000000e+00	
75%	0.000000e+00	2.100000e+01	0.000000e+00	1.500000e+01	0.000000e+00	
max	1.000000e+00	2.436000e+03	1.352000e+03	1.357000e+03	3.920000e+02	

  

	LateAircraftDelay
count	1.247488e+06
mean	2.529647e+01
std	4.205486e+01
min	0.000000e+00
25%	0.000000e+00
50%	8.000000e+00
75%	3.300000e+01
max	1.316000e+03

[8 rows x 25 columns]

Exploratory data analysis. Describe method brings us count, mean, median vs. values of the dataset.

```
[4]: df.shape
```

```
[4]: (1936758, 30)
```

Here is the number of rows and columns.

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1936758 entries, 0 to 1936757
Data columns (total 30 columns):
Unnamed: 0          int64
Year                int64
```

```

Month                int64
DayofMonth           int64
DayOfWeek            int64
DepTime              float64
CRSDepTime           int64
ArrTime              float64
CRSArrTime           int64
UniqueCarrier        object
FlightNum            int64
TailNum              object
ActualElapsedTime     float64
CRSElapsedTime       float64
AirTime              float64
ArrDelay             float64
DepDelay             float64
Origin               object
Dest                 object
Distance             int64
TaxiIn               float64
TaxiOut              float64
Cancelled            int64
CancellationCode      object
Diverted             int64
CarrierDelay         float64
WeatherDelay         float64
NASDelay             float64
SecurityDelay        float64
LateAircraftDelay    float64
dtypes: float64(14), int64(11), object(5)
memory usage: 443.3+ MB

```

Exploratory data analysis. Different data types in the dataset by columns.

```

[6]: Le_UniqueCarrier=LabelEncoder()
     Le_Origin=LabelEncoder()
     Le_Dest=LabelEncoder()
     Le_CancellationCode=LabelEncoder()

```

We have already detected the object type columns and now we convert those object-type columns into numerical values. Since machine learning algorithms cannot deal with strings.

```

[7]: df['UniqueCarrier_n'] = Le_UniqueCarrier.fit_transform(df['UniqueCarrier'])
     df['Origin_n'] = Le_Origin.fit_transform(df['Origin'])
     df['Dest_n'] = Le_Dest.fit_transform(df['Dest'])
     df['CancellationCode_n'] = Le_CancellationCode.
     ↪fit_transform(df['CancellationCode'])

```

Object-type columns are converted into numerical type columns. New numerical columns are inserted into our dataset.

```
[8]: df.shape
```

```
[8]: (1936758, 34)
```

4 brand new columns are added.

```
[9]: df.drop(['UniqueCarrier', 'Origin', 'Dest', 'CancellationCode', 'TailNum',  
↳ 'Unnamed: 0'], axis = 1, inplace=True)
```

4 Object-type columns, 1 extra index column and 1 other irrelevant (tail number) column are dropped.

```
[10]: df.shape
```

```
[10]: (1936758, 28)
```

The number of columns and rows are checked to see if adding and dropping are successful. So far so good.

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1936758 entries, 0 to 1936757  
Data columns (total 28 columns):  
Year                int64  
Month               int64  
DayOfMonth          int64  
DayOfWeek           int64  
DepTime             float64  
CRSDepTime          int64  
ArrTime             float64  
CRSArrTime          int64  
FlightNum           int64  
ActualElapsedTime   float64  
CRSElapsedTime      float64  
AirTime             float64  
ArrDelay            float64  
DepDelay            float64  
Distance            int64  
TaxiIn              float64  
TaxiOut             float64  
Cancelled           int64  
Diverted            int64  
CarrierDelay        float64  
WeatherDelay        float64  
NASDelay            float64  
SecurityDelay       float64  
LateAircraftDelay   float64  
UniqueCarrier_n     int64
```

```
Origin_n          int64
Dest_n            int64
CancellationCode_n int64
dtypes: float64(14), int64(14)
memory usage: 413.7 MB
```

The data type of each column is checked.

```
[12]: df.isnull().sum()
```

```
[12]: Year          0
      Month         0
      DayofMonth    0
      DayOfWeek     0
      DepTime       0
      CRSDepTime    0
      ArrTime       7110
      CRSArrTime    0
      FlightNum     0
      ActualElapsedTime 8387
      CRSElapsedTime 198
      AirTime       8387
      ArrDelay      8387
      DepDelay       0
      Distance      0
      TaxiIn        7110
      TaxiOut       455
      Cancelled     0
      Diverted      0
      CarrierDelay  689270
      WeatherDelay  689270
      NASDelay      689270
      SecurityDelay 689270
      LateAircraftDelay 689270
      UniqueCarrier_n 0
      Origin_n      0
      Dest_n        0
      CancellationCode_n 0
      dtype: int64
```

Null values are checked.

```
[13]: df.dropna(axis=0, how='any', inplace=True)
```

Rows that have null values are dropped.

```
[14]: df.isnull().sum()
```

```
[14]: Year          0
      Month         0
      DayOfMonth    0
      DayOfWeek     0
      DepTime       0
      CRSDepTime    0
      ArrTime       0
      CRSArrTime    0
      FlightNum     0
      ActualElapsedTime 0
      CRSElapsedTime 0
      AirTime       0
      ArrDelay      0
      DepDelay      0
      Distance      0
      TaxiIn        0
      TaxiOut       0
      Cancelled     0
      Diverted      0
      CarrierDelay  0
      WeatherDelay  0
      NASDelay      0
      SecurityDelay 0
      LateAircraftDelay 0
      UniqueCarrier_n 0
      Origin_n      0
      Dest_n        0
      CancellationCode_n 0
      dtype: int64
```

No more null values!

```
[15]: df.reset_index(drop=True, inplace=True)
```

We reset index to tidy up our data set.

```
[16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1247488 entries, 0 to 1247487
Data columns (total 28 columns):
Year          1247488 non-null int64
Month         1247488 non-null int64
DayOfMonth    1247488 non-null int64
DayOfWeek     1247488 non-null int64
DepTime       1247488 non-null float64
CRSDepTime    1247488 non-null int64
ArrTime       1247488 non-null float64
```

```

CRSArrTime      1247488 non-null int64
FlightNum       1247488 non-null int64
ActualElapsedTime 1247488 non-null float64
CRSElapsedTime  1247488 non-null float64
AirTime         1247488 non-null float64
ArrDelay        1247488 non-null float64
DepDelay        1247488 non-null float64
Distance        1247488 non-null int64
TaxiIn          1247488 non-null float64
TaxiOut         1247488 non-null float64
Cancelled       1247488 non-null int64
Diverted        1247488 non-null int64
CarrierDelay    1247488 non-null float64
WeatherDelay    1247488 non-null float64
NASDelay        1247488 non-null float64
SecurityDelay   1247488 non-null float64
LateAircraftDelay 1247488 non-null float64
UniqueCarrier_n 1247488 non-null int64
Origin_n        1247488 non-null int64
Dest_n          1247488 non-null int64
CancellationCode_n 1247488 non-null int64
dtypes: float64(14), int64(14)
memory usage: 266.5 MB

```

Dataset is controlled for last time and now we are ready to train our model.

```
[17]: y=df.loc[:, 'DepDelay']
```

Target variable is set.

```
[18]: X=df.loc[:, df.columns != 'DepDelay']
```

The rest of the columns will be used as predictors.

```
[19]: X_train, X_test, y_train, y_test= train_test_split(X,y, test_size=0.2,
↳random_state=12)
```

Train and test data are splitted. random\_state is used to get same result if code is implemented later again.

```
[20]: param_grid={"n_estimators":randint(1,9),
                  "max_depth": (2,7),
                  "max_features": randint(1,9),
                  "min_samples_leaf": randint(1,9)}
```

Hyperparameter tuning method is decided to choose parameters such as max\_depth, n\_estimators, max\_features, and min\_sample\_leaf.

```
[21]: RFReg=RandomForestRegressor(n_jobs=-1, random_state=42)
```

RandomForestRegressor method is used. `n_job=-1` helps computer to use full capacity CPU. `random_state` is again for later implementations of the same code.

```
[22]: RFReg_cv= RandomizedSearchCV(RFReg, param_grid, cv=5)
```

I decided to use `RandomizedSearchCV` rather than plain hyper parameter tuning (`GridSearchCV`) to make parameter choosing process shorter.

```
[23]: RFReg_cv.fit(X_train, y_train)
```

```
[23]: RandomizedSearchCV(cv=5, error_score='raise-deprecating',
                        estimator=RandomForestRegressor(bootstrap=True,
                                                         criterion='mse',
                                                         max_depth=None,
                                                         max_features='auto',
                                                         max_leaf_nodes=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         n_estimators='warn',
                                                         n_jobs=-1, oob_score=False,
                                                         random_state...
                        'max_features':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7f73676bbfd0>,
                        'min_samples_leaf':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7f73669baed0>,
                        'n_estimators':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7f73669ba5d0>},
                        pre_dispatch='2*n_jobs', random_state=None, refit=True,
                        return_train_score=False, scoring=None, verbose=0)
```

Training is implemented.

```
[24]: RFReg_cv.best_params_
```

```
[24]: {'max_depth': 7, 'max_features': 6, 'min_samples_leaf': 7, 'n_estimators': 5}
```

These parameters are chosen by the algorithm as the most optimized ones.

```
[25]: y_pred=RFReg_cv.predict(X_test)
```

Prediction is implemented on the test data.

```
[26]: mean_squared_error(y_test, y_pred)
```

```
[26]: 337.56074148431287
```



Calculation of MSE.

```
[27]: rmse=mean_squared_error(y_test, y_pred)**0.5
```

Calculation of average value error.

```
[28]: rmse
```

```
[28]: 18.372826170306865
```

It is pretty close to the value that I found by implementing linear regression, yet slightly worse than that.

```
[29]: RFReg_cv.score(X_test, y_test)
```

```
[29]: 0.9052527483794639
```

Calculation of R2

## The Theory Behind My Solution

1) In first assignment, I used a feature selection method to select most relevant features. However, in this assignment I didn't use feature selection. Because second assignment is still a regression problem, and although there are lots of different feature selection algorithms, same feature selection algorithm would apply to this problem. The reason is that both input and output data are numerical. So, I wanted to see the difference how feature selection algorithm affects the result as well as random forest method.

2) I converted string type columns into numeric type. Because ML algorithms cannot handle string type data. In the first assignment I dropped string type columns thinking that they are irrelevant, as a matter of fact, I was supposed to leave this decision(column eliminating decision) to feature selection algorithm. So, this is expected to improve the result of a model.

3) Some resources indicate that we do not need to drop null values in decision tree/random forest model because the process itself already handles it, yet I did drop null values.

4) The biggest factor that may differ my model from others is hyper parameter tuning. In decision tree/random forest method, we need to detect the depth of each tree, the number of trees and the number of leaf in each tree vs. If we (rather than algorithm) decide all these details, the power of model will be limited. However, if we leave this decision to an unsupervised algorithm(hyper parameter tuning algorithm) then algorithm will find the most optimized parameters of algorithm on its own. Here we have two options: One is gridsearch which try each and every single parameter in the algorithm. This one takes way too long to process. The other one is randomized search. This one takes relatively shorter because it doesnt try every single parameter. It chooses a few parameters randomly and tries to find most optimized ones from randomly selected parameters. I used the latter as the computation capacity of my Personal Computer is limited.

5) I want to mention "n\_jobs=-1" more particularly which does pretty good job to use whole computational capacity of CPU, thanks to which it takes shorter to train your data.

6) The result : R2 score is slightly worse than my previous model which is linear regression. (90.5% vs 92.6%) The reason might be feature selection algorithm or even maybe random\_state number. As this is not a classification problem evaluation method cannot be confusion matrix. RMSE, MSE and accuracy score are only ways to measure the performance of our model.